

THE HARNESS

Same model. Very different results.

CLAUDE CODE · CODEX · DEEPSEEK HARNESS

8 slides · what actually decides agent performance

THE MODEL IS NOT THE AGENT.

The harness decides what the model sees,
which tools it can call, how state persists,
when it retries, and when it stops.

Same model, four harnesses.
Different scores. Different cost.

The wrapper shapes the inference, not just the output.

ONE MODEL, THREE SCORES

DeepSeek V4 Pro through each harness:

Pi	21 / 30
DeepSeek Harness ..	20 / 30
Claude Code	19 / 30

Tokens per task: 88k | 650k | 925k

Success rate and token cost move in opposite directions.

AGENT-CENTRIC

- 26 lifecycle hooks: inject logic at every step
- CLAUDE.md, skills, subagents, MCP, plugins
- 1M-token context; fastest median time in test
- broadest ecosystem of the three
- closed source; optimised for Claude models

Make one agent stronger. Product, not toolbox.

EXECUTION-CENTRIC

- kernel-level sandbox: Seatbelt, Landlock
- boundaries the model cannot talk its way past
- AGENTS.md spec; 272K context, to 1.05M
- Codex Cloud: up to 6 parallel threads
- safety leader for untrusted code

Run real systems without losing control of them.

RUNTIME-CENTRIC

- Cordis microkernel: everything is a plugin
- models, tools, sessions, sandboxes, loop, UI
- 40+ model providers, model-agnostic by design
- trajectory view: every action traced to its plugin
- can call Claude Code / Codex as sub-agents
- v0.1 preview: breaking changes expected

A frame you build on, not a product you adopt.

THE TOKEN TAX

Framework overhead, input tokens per task:

- OpenCode ~12k
- Codex CLI ~18k
- Claude Code ~22k
- DeepSeek Harness .. ~95k

Confirmed bug: duplicate CLAUDE.md + AGENTS.md
injection adds 30k+ tokens on its own.

Flexibility is paid for in context, every turn.

WHO OWNS THE DECISIONS?

Model. Runtime. Policy. Human.

That split is the real difference.

Run your own 10-20 task benchmark before you commit.